

date: See Website.  
TA in charge: See Website.

**Please note:** You should not publish your lab solutions in any publicly accessible site such as GitHub.

## Lab Q&A

We encourage you to ask questions on course's Piazza forum. If no help provided on Piazza forum then try to email TA in charge. E-mails regarding this lab (such as administrative issues) should be sent with the subject "OSE, lab1".

## Introduction

This lab is split into three parts. The first part concentrates on getting familiarized with x86 assembly language, the QEMU x86 emulator, and the PC's power-on bootstrap procedure. The second part examines the boot loader for our kernel, which resides in the boot directory of the lab tree. Finally, the third part delves into the initial template for our kernel itself, named JOS, which resides in the kernel directory.

## Machine Setup

You are required to run Linux based distribution on your PC, namely, Ubuntu Desktop 14.04 for i386 architecture. However, you don't have to wipe your machine and install Ubuntu. You can use dual-boot option, or you can run it as a virtual machine.

**Note:** You can use physical machine instead of VM and skip the first step. If you're advanced Linux user and already running the required Linux (see below), you can even skip first two steps. However, do this on your own risk, we wouldn't be responsible if you suddenly loose all your data. (If you're experienced Linux user, you can substitute steps 1&2 with chroot creation.

### Step 1. VM Setup

Any decent virtual machine software will suit our needs. For example, free and open source [VirtualBox](#), or commercial-only [VmWare Workstation](#). The later is installed on some PC farm stations and can be [obtained through CS2.4 faculty](#) in no charge. Create new VM for our Linux and ensure it has access to the Internet (we'll need it later).

**Step 2. Linux installation**  
Download [Ubuntu 14.04.2 LTS](#), [Ubuntu 64-bit x86\\_64](#) and install it inside VM. Note, that other versions of Ubuntu, as well as other Linux distributions **aren't suitable for us** (Ubuntu 14.04.x is considered the same as 14.04). Feel free to install your favorite code editors or other software inside this "virtual" Ubuntu.

### Step 3. Additional Software Installation

Login into your Ubuntu, open terminal and run (hereinafter \$ denotes a shell prompt):

```
$ sudo apt-get update
[sudo] password for user: <enter your user password (not echoed to the terminal)>
<lots of lines>
$ sudo apt-get -y install qemu-gui git build-essential gcc-multilib
<lots of lines here>
$ sudo ln -s /usr/bin/qemu-system-i386 /usr/bin/qemu
$ echo "set auto-load safe-path /" >> ~/.gnupg/
$ cd lab
```

In order to save you some typing you can open this page inside VM (there is a Firefox browser installed by default) and copy-paste long commands into the terminal (of course, you already know that it's often sufficient to just select the text by clicking with the left mouse button on the mouse to paste it, in order to press Ctrl-C/Ctrl-V).

### Source Code Setup

The files you will need for this and subsequent lab assignments in this course are distributed using the [Git](#) version control system. To learn more about Git, take a look at the [Git User's manual](#), or, if you are already familiar with other version control systems, you may find this [CS-oriented overview of Git](#) useful.

The URL for the course Git repository is <http://www.cs.techion.ac.il/~cs236376/jos-gis>. To install the files in your machine, you need to *clone* the course repository, by running the commands below.

```
$ git clone http://www.cs.tchion.ac.il/~cs236376/jos-gis lab
Initialized empty git repository in /home/user/lab/.git/
$ cd lab
```

Git allows you to keep track of the changes you make to the code. For example, if you are finished with one of the exercises, and want to checkpoint your progress, you can *commit* your changes by running:

```
$ git commit -am "my solution for lab1 exercises"
Created commit 0a21251: my solution for lab1 exercises
1 files changed, 1 insertion(+), 0 deletion(-)
```

Git is a distributed version control system, unlike SVN or CVS, which means all your commits will be done to a local copy of the repository and are only saved on your own computer.

You can keep track of your changes by using the `git diff` command. Running `git diff` will display the changes to your code since your last commit, and `git diff origin:lab1` will display the changes relative to the initial code supplied for this lab. Here, `origin:lab1` is the name of the git branch with the initial code you downloaded from our server for this assignment.

### Questionary

In addition to writing few lines of code, you need to answer several questions throughout the lab. A [textual template](#) called `lab1-questionary.txt` contains the required questions formulated in easy to check manner.

Download it to lab1 source directory (the content of the file is shorter and easier to navigate than more recent manuals but describes all of the x86 processor features that we will make use of in the course and the full, latest and greatest `lab1-questionary.txt` to fill the answer. Some questions appearing here are for self-check and don't require answer submission (so they are absent from `lab1-questionary.txt`).

The challenge for this lab is completely optional and will not be checked.

### Hand-In (Submission) Procedure

When you are ready to hand in your lab (including the filled `lab1-questionary.txt`), run `make handin` in the source directory. This will make a tar file, which you can then submit via [webcourse site](#). You can list the contents of the tar file with `tar -tvf lab1-handin.tar.gz` or unpack it (in another directory) with `tar -xzf lab1-handin.tar.gz`.

We will be grading your solutions with a grading program. You can run `make grade` to test your code with the grading program (no test for the questionary is provided).

## Part 1: PC Bootstrapping

The purpose of the first exercise is to introduce you to x86 assembly language and the PC bootstrapping process, and to get you started with QEMU and QEMU/GDB debugging. You will not have to write any code for this part of the lab, but you should go through it anyway for your own understanding and be prepared to answer the questions posed below.

### Getting Started with x86 assembly

If you are not already familiar with x86 assembly language, you will quickly become familiar with it during this course! The [PC Assembly Language Book](#) is an excellent place to start. Hopefully, the book contains mixture of new and old material for you.

**Warning:** Unfortunately the examples in the book are written for the NASM assembler, whereas we will be using the GNU assembler. NASM uses the so-called *Intel* syntax while GNU uses the *AT&T* syntax. While semantically equivalent, an assembly file will differ quite a bit, and in fact superficially, depending on which syntax is used. Luckily the conversion between the two is pretty simple, and is covered in [Brennan's Guide to Inline Assembly](#).

Certainly the definitive reference for x86 assembly language programming is Intel's instruction set architecture reference, which you can find on [the source reference page](#) in two flavors: an HTML edition of the [80386 Programmer's Reference Manual](#), which is much shorter and easier to navigate than more recent manuals but describes all of the x86 processor features that we will make use of in the course and the full, latest and greatest [IA-32 Intel Architecture Software Developer's Manuals](#) from Intel, covering all the features of the most recent processors that we won't need in class but you may be interested in learning about. An equivalent (and often friendlier) set of manuals is [available from AMD](#). Save the Intel/AMD architecture manuals for later or use them for reference when you want to look up the definitive explanation of a particular processor feature or instruction.

### Simulating the x86

Instead of developing the operating system on a real personal computer (PC), we use a program that faithfully emulates a complete PC: the code you write for the emulator will boot on a real PC too. Using an emulator simplifies debugging; you can, for example, set break points inside the emulated x86, which is difficult to do with the silicon version of an x86.

In the course we will use the [QEMU Emulator](#), a modern and relatively fast emulator. While QEMU's built-in monitor provides only limited debugging support, QEMU can act as a remote debugging target for the [GNU debugger](#) (GDB), which will use this in this lab to step through the early boot process.

To get started, external the lab's files as described above in "Software Setup", then type `make` in the lab directory to build the minimal boot loader and kernel you will start with. (It's a little generous to call the code we're running here a "kernel", but we'll flesh it out throughout the semester.)

```
$ cd lab
$ make
+ cc kern/entry.S
+ cc kern/init.c
+ cc kern/console.c
+ cc kern/main.c
+ cc kern/print.c
+ cc lib/print.c
+ cc lib/readline.c
+ cc lib/string.c
+ ld obj/kern/kernel
+ cc boot/boot.c
+ cc obj/boot/main.o
+ ld boot/boot
+ cc boot/boot.o -o boot/boot -s 512
+ cc obj/kern/kernel.o
```

Now you're ready to run QEMU, supplying the file `obj/kern/kernel.img`, created above, as the contents of the emulated PC's "virtual hard disk." This hard disk image contains both our boot loader (`obj/boot/boot`) and our kernel (`obj/kernel`).

This executes QEMU with the options required to set the hard disk and direct serial port output to the terminal. Some text should appear in the QEMU window:

```
Booting from Hard Disk...
DOS bootsector 0x00000000
entering test_backtrace 5
entering test_backtrace 4
entering test_backtrace 3
entering test_backtrace 2
entering test_backtrace 0
leaving test_backtrace 0
leaving test_backtrace 2
leaving test_backtrace 3
leaving test_backtrace 4
leaving test_backtrace 5
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
```

Everything after "booting from Hard Disk..." was printed by our skeletal JOS kernel; the `k>` is the prompt printed by the small *monitor*, or interactive control program, that we've included in the kernel. These lines printed by the kernel will also appear in the regular shell window from which you run QEMU. This is because for testing and lab grading purposes we have set up the JOS kernel to write its console output not only to the virtual VGA display (as seen in the QEMU window), but also to the simulated PC's virtual serial port, which QEMU in turn outputs to its own standard output. Likewise, the JOS kernel will take input from both the keyboard and the serial port, so you can give it commands in either the VGA display window or the terminal running QEMU. Alternatively, you can use the serial console without the virtual VGA by running `make qemu-nox`.

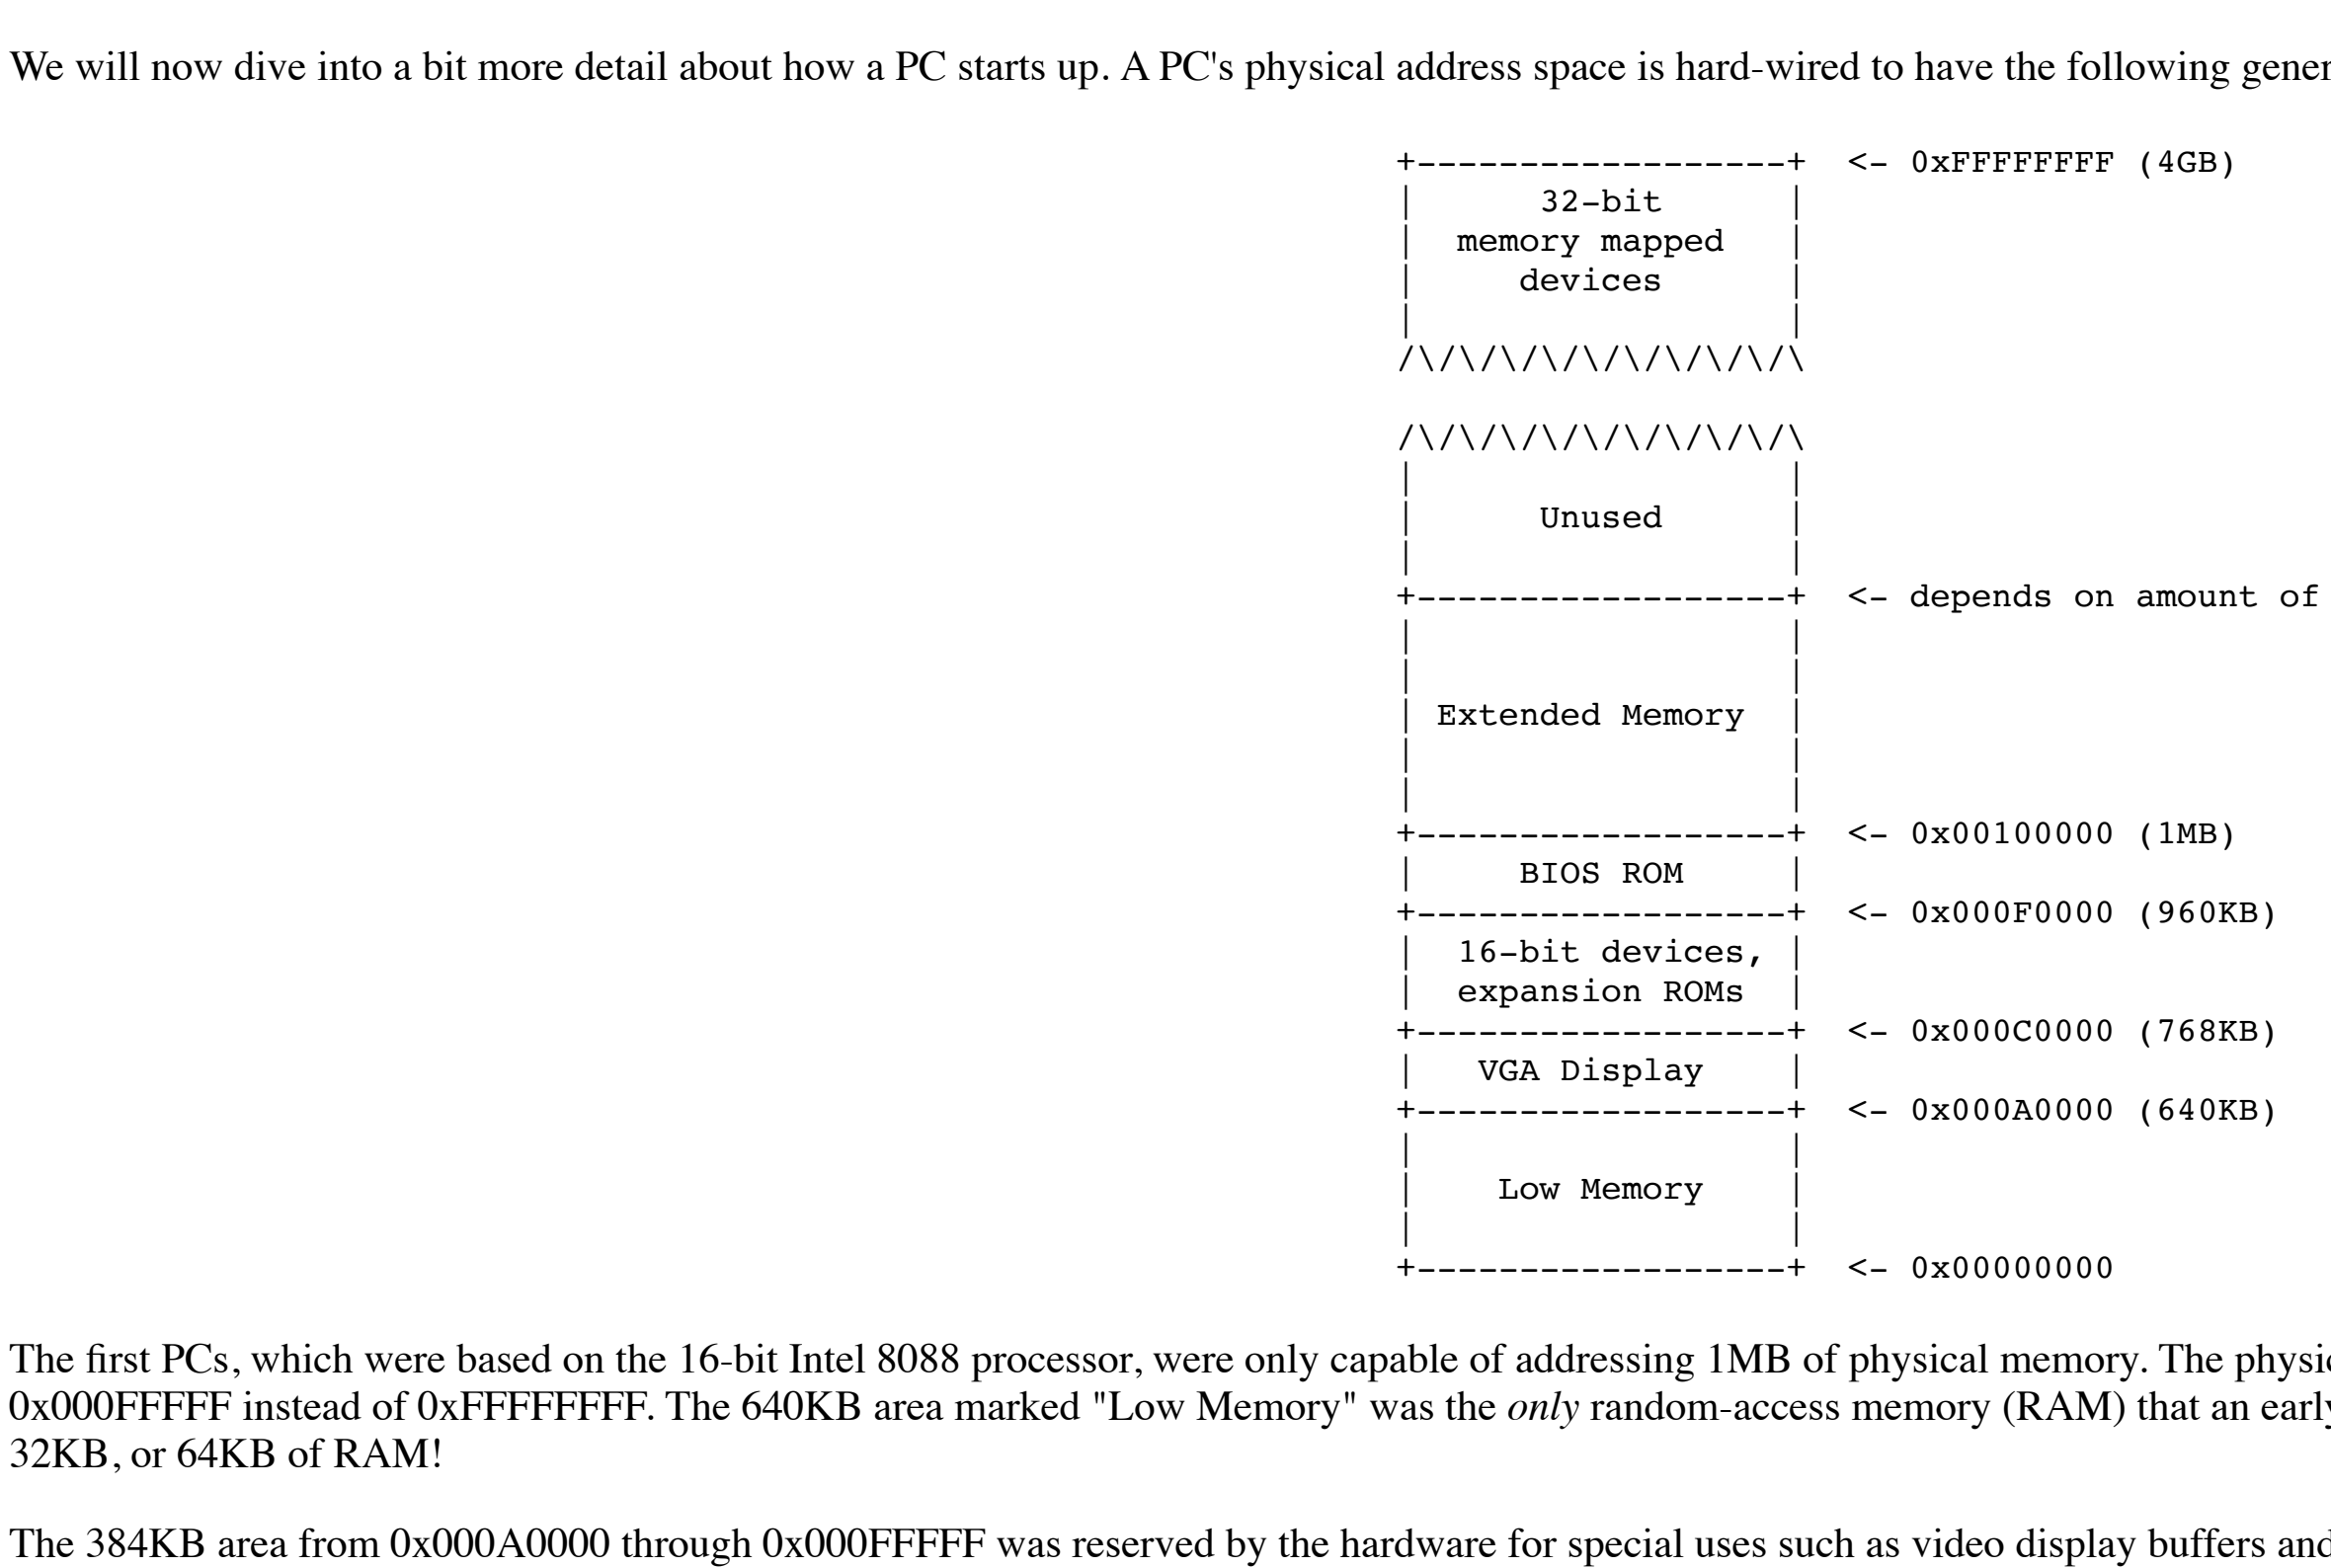
There are only two commands you can give to the kernel without the VGA display, and they are `help` and `kernelinfo`.

```
k> help
help - display this list of commands
kernelinfo - display information about the kernel
k> kernelinfo
Special kernel symbols:
entry 00100000 (virt) 00100000 (phys)
start 00101075 (virt) 00101075 (phys)
starta 00102000 (virt) 00102000 (phys)
end 00112600 (virt) 00112600 (phys)
enda 00112600 (virt) 00112600 (phys)
kernel executable memory footprint: 19136
```

The help command is obvious, and we will shortly discuss the meaning of what the `kernelinfo` command prints. Although simple, it's important to note that this kernel monitor is running "directly" on the "raw (virtual) hardware" of the simulated PC. This means that you should be able to copy the contents of `obj/kern/kernel.img` onto the first few sectors of a *real* hard disk, insert that hard disk into a real PC, turn it on, and see exactly the same thing on the PC's real screen as you did above in the QEMU window. (We don't recommend you do this on a real machine with useful information on its hard disk, though, because copying `kernel.img` onto the beginning of its hard disk will trash the master boot record and the beginning of the first partition, effectively causing everything previously on the hard disk to be lost!)

### The PC's Physical Address Space

We will now dive into a bit more detail about how a PC starts up. A PC's physical address space is hard-wired to have the following general layout:



The first PCs, which were based on the 16-bit Intel 8088 processor, were only capable of addressing 1MB of physical memory. The physical address space of an early PC would therefore start at 0x00000000 but end at 0x000FFFFF instead of 0xFFFFFFFF. The 640KB area marked "Low Memory" was the *only* random-access memory (RAM) that an early PC could use; in fact the very earliest PCs only could be configured with 16KB, 32KB, or 64KB of RAM!

The 384KB area from 0x00A0000 through 0x00FFFFFFF was reserved by the hardware for special uses such as video display buffers and firmware held in non-volatile memory. The most important part of this reserved area is the Basic Input/Output System (BIOS), which occupies the 64KB region from 0x00000000 through 0x000FFFFF. In early PCs the BIOS was held in read-only memory (ROM), but current PCs store the BIOS in erasable (and often rewritable) non-volatile memory, which is much shorter and easier to navigate than more recent manuals but describes all of the x86 processor features that we will make use of in the course and the full, latest and greatest [IA-32 Intel Architecture Software Developer's Manuals](#) from Intel, covering all the features of the most recent processors that we won't need in class but you may be interested in learning about. An equivalent (and often friendlier) set of manuals is [available from AMD](#). Save the Intel/AMD architecture manuals for later or use them for reference when you want to look up the definitive explanation of a particular processor feature or instruction.

When Intel finally "broke the one megabyte barrier" with the 80286 and 80386 processors, which supported 16MB and 4GB physical address spaces respectively, the PC architects nevertheless preserved the original layout for the low 1MB of physical address space in order to ensure backward compatibility with existing software. Modern PCs therefore have a "hole" in physical memory from 0x00A0000 to 0x00FFFFFFF, dividing RAM into "low" or "conventional memory" (the first 640KB) and "extended memory" (everything else). In addition, some space at the very top of the PC's 32-bit physical address space, above all physical RAM, is now commonly reserved by the BIOS for use by 32-bit PCI devices.

Recent x86 processors can support more than 4GB of physical RAM, so RAM can extend further above 0xFFFFFFFF. In this case the BIOS must arrange to leave a *second* hole in the system's RAM at the top of the 32-bit addressable region to leave room for the first 32-bit devices to be mapped. Because of design limitations JOS will use only the first 256MB of a PC's physical memory anyway, so for now we will pretend that all PCs have only 32-bit physical address space. But dealing with complicated physical address spaces and limitations aspects of hardware organization that evolved over many years is one of the important practical challenges of OS development.

### The ROM BIOS

In this portion of the lab, you'll use QEMU's debugging facilities to investigate how an IA-32 compatible computer boots.

Open two terminal windows. In one, enter `make qemu-gdb` (or `make qemu-nox-gdb`). This starts up QEMU, but QEMU stops just before the processor executes the first instruction and waits for a debugging connection from GDB. In the second terminal, from the same directory you ran `make`, run `gdb`. You should see something like this:

```
$ gdb
GNU gdb (Ubuntu 7.7.1-1ubuntu5~14.04.2) 7.7.1
Copyright (c) 2014 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type 'show copying'
and 'show warranty' for details.
GDB was configured as 'x86_64-linux-gnu'.
Type 'show configuration' for configuration details.
For help, type 'help'.
> file obj/kern/kernel
obj/kern/kernel: ELF 32-bit LSB executable, Intel 80386, version 1.0, dynamically linked,
  <http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type 'help'.
> target remote localhost:26000
warning: A handler for the OS ABI "GNU/Linux" is not built into this
configuration of GDB.  Attempting to continue with the default i386 settings.
The target architecture is assumed to be i386
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
0xffff: 0x71
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
(gdb)
```

We provided a `.gdbinit` file that set up GDB to debug the 16-bit code used during early boot and directed it to attach to the listening QEMU.

The following line:

```
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
```

is GDB's disassembly of the first instruction to be executed. From this output you can conclude a few things:

- The IBM PC starts executing at physical address 0x0000ffh, which is at the very top of the 64KB area reserved for the BIOS.
- The PC starts executing with `cs = 0xf000` and `ip = 0xffff`.
- The first instruction to be executed is a `jmp` instruction, which jumps to the segmented address `cs = 0xf000` and `ip = 0x005b`.

Why does QEMU start like that? This is how Intel designed the 8088 processor, which IBM used in their original PC. Because the BIOS in a PC is "hard-wired" to start at 0x00000000.0x0000ffh, this is where memory at `0x00000000` (Note that `0x00000000` is the address of the first instruction to be executed, and that after the instruction offsets are 32 bits instead of 16).

Protected mode is described briefly in sections 1.2.7 and 1.2.8 of [PC Assembly Language](#), and in great detail in the Intel architecture manuals. At this point you only have to understand that translation of segmented addresses (so-called *segment:offset* pairs) into physical addresses happens differently in protected mode, and that after the translation offsets are 32 bits instead of 16.

Second, the boot loader reads the kernel from the hard disk by directly accessing the IDE disk device registers via the x86's special I/O instructions. If you would like to understand better what the particular I/O instructions here mean, check out the "IDE hard drive controller" section on [the source reference page](#). You will not need to learn much about programming specific devices in this class: writing device drivers is in practice a very important part of OS development, but from a conceptual or architectural viewpoint it is also one of the least interesting.

After you understand the boot loader source code, look at the file `obj/boot/boot.asm`. This file is a disassembly of the boot loader that our GNU/makefile creates after compiling the boot loader. This disassembly file makes it easy to see exactly where in physical memory all of the boot loader code resides, and makes it easier to track what's happening while stepping through the instructions `gdb` at a time.

By the virtual memory hardware to physical addresses, `entry` points translates virtual addresses in the range 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000, as well as virtual addresses 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000. Any virtual address that is not in one of these two ranges will cause a hardware interrupt (called a page fault) which aren't able to handle just yet.

You can set address breakpoints in GDB with the `b` command. For example, `b *0x005b` sets a breakpoint at address 0x00005b. Once at a breakpoint, you can continue execution using the `c` and `s` commands: `c` causes QEMU to continue execution until the next breakpoint (or until you press Ctrl-C in GDB), and `s` steps through the instructions `gdb` at a time.

To examine instructions in memory (besides the immediate next one to be executed, which GDB prints automatically), you use the `x/N` command. This command has the syntax `x/N <addr>`, where `N` is the number of consecutive instructions to disassemble and `ADDR` is the memory address at which to start disassembling.

The following line:

```
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
```

is GDB's disassembly of the first instruction to be executed. From this output you can conclude a few things:

- The IBM PC starts executing at physical address 0x0000ffh, which is at the very top of the 64KB area reserved for the BIOS.
- The PC starts executing with `cs = 0xf000` and `ip = 0xffff`.
- The first instruction to be executed is a `jmp` instruction, which jumps to the segmented address `cs = 0xf000` and `ip = 0x005b`.

Why does QEMU start like that? This is how Intel designed the 8088 processor, which IBM used in their original PC. Because the BIOS in a PC is "hard-wired" to start at 0x00000000.0x0000ffh, this is where memory at `0x00000000` (Note that `0x00000000` is the address of the first instruction to be executed, and that after the instruction offsets are 32 bits instead of 16).

Protected mode is described briefly in sections 1.2.7 and 1.2.8 of [PC Assembly Language](#), and in great detail in the Intel architecture manuals. At this point you only have to understand that translation of segmented addresses (so-called *segment:offset* pairs) into physical addresses happens differently in protected mode, and that after the translation offsets are 32 bits instead of 16.

Second, the boot loader reads the kernel from the hard disk by directly accessing the IDE disk device registers via the x86's special I/O instructions. If you would like to understand better what the particular I/O instructions here mean, check out the "IDE hard drive controller" section on [the source reference page](#). You will not need to learn much about programming specific devices in this class: writing device drivers is in practice a very important part of OS development, but from a conceptual or architectural viewpoint it is also one of the least interesting.

After you understand the boot loader source code, look at the file `obj/boot/boot.asm`. This file is a disassembly of the boot loader that our GNU/makefile creates after compiling the boot loader. This disassembly file makes it easy to see exactly where in physical memory all of the boot loader code resides, and makes it easier to track what's happening while stepping through the instructions `gdb` at a time.

By the virtual memory hardware to physical addresses, `entry` points translates virtual addresses in the range 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000, as well as virtual addresses 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000. Any virtual address that is not in one of these two ranges will cause a hardware interrupt (called a page fault) which aren't able to handle just yet.

You can set address breakpoints in GDB with the `b` command. For example, `b *0x005b` sets a breakpoint at address 0x00005b. Once at a breakpoint, you can continue execution using the `c` and `s` commands: `c` causes QEMU to continue execution until the next breakpoint (or until you press Ctrl-C in GDB), and `s` steps through the instructions `gdb` at a time.

To examine instructions in memory (besides the immediate next one to be executed, which GDB prints automatically), you use the `x/N` command. This command has the syntax `x/N <addr>`, where `N` is the number of consecutive instructions to disassemble and `ADDR` is the memory address at which to start disassembling.

The following line:

```
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
```

is GDB's disassembly of the first instruction to be executed. From this output you can conclude a few things:

- The IBM PC starts executing at physical address 0x0000ffh, which is at the very top of the 64KB area reserved for the BIOS.
- The PC starts executing with `cs = 0xf000` and `ip = 0xffff`.
- The first instruction to be executed is a `jmp` instruction, which jumps to the segmented address `cs = 0xf000` and `ip = 0x005b`.

Why does QEMU start like that? This is how Intel designed the 8088 processor, which IBM used in their original PC. Because the BIOS in a PC is "hard-wired" to start at 0x00000000.0x0000ffh, this is where memory at `0x00000000` (Note that `0x00000000` is the address of the first instruction to be executed, and that after the instruction offsets are 32 bits instead of 16).

Protected mode is described briefly in sections 1.2.7 and 1.2.8 of [PC Assembly Language](#), and in great detail in the Intel architecture manuals. At this point you only have to understand that translation of segmented addresses (so-called *segment:offset* pairs) into physical addresses happens differently in protected mode, and that after the translation offsets are 32 bits instead of 16.

Second, the boot loader reads the kernel from the hard disk by directly accessing the IDE disk device registers via the x86's special I/O instructions. If you would like to understand better what the particular I/O instructions here mean, check out the "IDE hard drive controller" section on [the source reference page](#). You will not need to learn much about programming specific devices in this class: writing device drivers is in practice a very important part of OS development, but from a conceptual or architectural viewpoint it is also one of the least interesting.

After you understand the boot loader source code, look at the file `obj/boot/boot.asm`. This file is a disassembly of the boot loader that our GNU/makefile creates after compiling the boot loader. This disassembly file makes it easy to see exactly where in physical memory all of the boot loader code resides, and makes it easier to track what's happening while stepping through the instructions `gdb` at a time.

By the virtual memory hardware to physical addresses, `entry` points translates virtual addresses in the range 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000, as well as virtual addresses 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000. Any virtual address that is not in one of these two ranges will cause a hardware interrupt (called a page fault) which aren't able to handle just yet.

You can set address breakpoints in GDB with the `b` command. For example, `b *0x005b` sets a breakpoint at address 0x00005b. Once at a breakpoint, you can continue execution using the `c` and `s` commands: `c` causes QEMU to continue execution until the next breakpoint (or until you press Ctrl-C in GDB), and `s` steps through the instructions `gdb` at a time.

To examine instructions in memory (besides the immediate next one to be executed, which GDB prints automatically), you use the `x/N` command. This command has the syntax `x/N <addr>`, where `N` is the number of consecutive instructions to disassemble and `ADDR` is the memory address at which to start disassembling.

The following line:

```
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
```

is GDB's disassembly of the first instruction to be executed. From this output you can conclude a few things:

- The IBM PC starts executing at physical address 0x0000ffh, which is at the very top of the 64KB area reserved for the BIOS.
- The PC starts executing with `cs = 0xf000` and `ip = 0xffff`.
- The first instruction to be executed is a `jmp` instruction, which jumps to the segmented address `cs = 0xf000` and `ip = 0x005b`.

Why does QEMU start like that? This is how Intel designed the 8088 processor, which IBM used in their original PC. Because the BIOS in a PC is "hard-wired" to start at 0x00000000.0x0000ffh, this is where memory at `0x00000000` (Note that `0x00000000` is the address of the first instruction to be executed, and that after the instruction offsets are 32 bits instead of 16).

Protected mode is described briefly in sections 1.2.7 and 1.2.8 of [PC Assembly Language](#), and in great detail in the Intel architecture manuals. At this point you only have to understand that translation of segmented addresses (so-called *segment:offset* pairs) into physical addresses happens differently in protected mode, and that after the translation offsets are 32 bits instead of 16.

Second, the boot loader reads the kernel from the hard disk by directly accessing the IDE disk device registers via the x86's special I/O instructions. If you would like to understand better what the particular I/O instructions here mean, check out the "IDE hard drive controller" section on [the source reference page](#). You will not need to learn much about programming specific devices in this class: writing device drivers is in practice a very important part of OS development, but from a conceptual or architectural viewpoint it is also one of the least interesting.

After you understand the boot loader source code, look at the file `obj/boot/boot.asm`. This file is a disassembly of the boot loader that our GNU/makefile creates after compiling the boot loader. This disassembly file makes it easy to see exactly where in physical memory all of the boot loader code resides, and makes it easier to track what's happening while stepping through the instructions `gdb` at a time.

By the virtual memory hardware to physical addresses, `entry` points translates virtual addresses in the range 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000, as well as virtual addresses 0x00000000 through 0x00000000 to physical addresses 0x00000000 through 0x00000000. Any virtual address that is not in one of these two ranges will cause a hardware interrupt (called a page fault) which aren't able to handle just yet.

You can set address breakpoints in GDB with the `b` command. For example, `b *0x005b` sets a breakpoint at address 0x00005b. Once at a breakpoint, you can continue execution using the `c` and `s` commands: `c` causes QEMU to continue execution until the next breakpoint (or until you press Ctrl-C in GDB), and `s` steps through the instructions `gdb` at a time.

To examine instructions in memory (besides the immediate next one to be executed, which GDB prints automatically), you use the `x/N` command. This command has the syntax `x/N <addr>`, where `N` is the number of consecutive instructions to disassemble and `ADDR` is the memory address at which to start disassembling.

The following line:

```
[<0000ffff>] 0xffff: jmp 0x0000,0x00005b
```

is GDB's disassembly of the first instruction to be executed. From this output you can conclude a few things:

- The IBM PC starts executing at physical address 0x0000ffh, which is at the very top of the 64KB area reserved for the BIOS.
- The PC starts executing with `cs = 0xf000` and `ip = 0xffff`.
- The first instruction to be executed is a `jmp` instruction, which jumps to the segmented address `cs = 0xf000` and `ip = 0x005b`.

Why does QEMU start like that? This is how Intel designed the 8088 processor, which IBM used in their original PC. Because the BIOS in a PC is "hard-wired" to start at 0x00000000.0x0000ffh, this is where memory at `0x00000000` (Note that `0x00000000` is the address of the first instruction to be executed, and that after the instruction offsets are 32 bits instead of 16).