

263676 OSE Lab 3: User Environments

Due date: See Webcourse.
TA in charge: See Webcourse

Please note: You should not publish your lab solutions in any publicly accessible site such as github.

Lab Q&A

We encourage you to ask questions on course's Piazza forum. If no help provided on Piazza forum then try to email TA in charge. E-mails regarding this lab (such as administrative issues) should be sent with the subject "OSE, Lab3".

Introduction

In this lab you will implement the basic kernel facilities required to get a protected user-mode environment (i.e., "process") running. You will enhance the JOS kernel to set up the data structures to keep track of user environments, create a single user environment, load a program image into it, and start it running. You will also make the JOS kernel capable of handling any system calls the user environment makes and handling any other exceptions it causes.

Note: In this lab, the terms *environment* and *process* are interchangeable – both refer to an abstraction that allows you to run a program. We introduce the term "environment" instead of the traditional term "process" in order to stress the point that JOS environments and UNIX processes provide different interfaces, and do not provide the same semantics.

Getting Started

Use `git` to commit your changes after your Lab 2 submission (if any), fetch the latest version of the course repository, and then create a local branch called `lab3` based on our `lab3` branch, `origin/lab3`:

```
$ git commit -am "changes to lab2 after handin"
Created commit 734fab7: changes to lab2 after handin
 4 files changed, 42 insertions(+), 9 deletions(-)
$ git pull
Already up-to-date.
$ git checkout -b lab3 origin/lab3
Branch lab3 set up to track remote branch refs/remotes/origin/lab3.
Switched to a new branch "lab3"
$ git merge lab2
Merge made by recursive.
 kern/panic.c | 42 ++++++++
 4 files changed, 42 insertions(+), 9 deletions(-)
$
```

Lab 3 contains a number of new source files, which you should browse:

```
inc/ env.h      Public definitions for user-mode environments
trap.h         Public definitions for trap handling
syscalls.h     Public definitions for system calls from user environments to the kernel
lib.h          Public definitions for the user-mode support library
kern/ env.h     Kernel-private definitions for user-mode environments
env.c          Kernel code implementing user-mode environments
trap.h         Kernel-private trap handling definitions
trap.c         Trap handling code
trapentry.S    Assembly-language trap handler entry-points
syscalls.h     Kernel-private definitions for system call handling
syscall.c      System call implementation code
lib/ makefrag  Makefile fragment to build user-mode library, obj/lib/libuser.a
entry.S        Assembly-language entry-point for user environments
libmain.c      User-mode library setup code called from user.S
syscall.c       User-mode system call stub functions
console.c      User-mode implementations of putchar and getchar, providing console I/O
exit.c         User-mode implementation of exit
panic.c        User-mode implementation of panic
user/ *        Various test programs to check kernel lab 3 code
```

In addition, a number of the source files we handed out for lab2 are modified in lab3. To see the differences, you can type:

```
$ git diff lab2
```

You may also want to take another look at the [lab3tools guide](#), as it includes information on debugging user code that becomes relevant in this lab.

Lab Requirements

As in lab 2, you will need to do all of the regular exercises described in the lab and at *least* one challenge problem. Write brief answers to the questions posed in the lab and a one or two paragraph description of what you did to solve your chosen challenge problem in the [questionary](#). (If you implement more than one challenge problem, you only need to describe one of them in the write-up.)

Hand-In Procedure

When you are ready to hand in your lab (including the filled [lab3-questionary.txt](#)), run `make handin` in the source directory. This will make a tar file for you, which you can then submit via [webcourse site](#). You can list the contents of the tar file with `tar -tvzf lab3-handin.tar.gz` or unpack it (in another directory) with `tar -xzf lab3-handin.tar.gz`.

As before, we will be grading your solutions with a grading program. You can run `make grade` in the lab directory to test your kernel with the grading program (no test for the environment is provided). The grading program may rely on some in-kernel code for the check. Needless to say, that altering this code or otherwise deceiving automatic testing is considered severe cheating.

Inline Assembly

In this lab you may find GCC's inline assembly language feature useful, although it is also possible to complete the lab without using it. At the very least, you will need to be able to understand the fragments of inline assembly language ("asm" statements) that already exist in the source code we gave you. You can find several sources of information on GCC's inline assembly language on the class [reference materials](#) page.

Part A: User Environments and Exception Handling

The new include file `inc/env.h` contains basic definitions for user environments in JOS. Read it now. The kernel uses the `env` data structure to keep track of each user environment. In this lab you will initially create just one environment, but you will need to design the JOS kernel to support multiple environments; lab 4 will take advantage of this feature by allowing a user environment to *fork* other environments.

As you can see in `kern/env.c`, the kernel maintains three main global variables pertaining to environments:

```
struct Env envs = NULL; // All environments
struct Env *current = NULL; // The current env
static struct Env *env_free_list; // Free environment list
```

Once JOS gets up and running, the `env` pointer points to an array of `Env` structures representing all the environments in the system. In our design, the JOS kernel will support a maximum of `ENV` simultaneously active environments, even though there will typically be far fewer running environments at any given time. (`ENV` is a constant *defined* in `inc/env.h`.) Once it is allocated, the `envs` array will contain a single instance of the `Env` data structure for each of the `ENV` possible environments.

The JOS kernel keeps all of the inactive `Env` structures on the `env_free_list`. This design allows easy allocation and deallocation of environments, as they merely have to be added to or removed from the free list.

The kernel uses the `curenv` symbol to keep track of the *currently executing* environment at any given time. During boot up, before the first environment is run, `curenv` is initially set to `NULL`.

Environment State

The `Env` structure is defined in `inc/env.h` as follows (although more fields will be added in future labs):

```
struct Env {
  struct Trapframe env_tf; // Saved registers
  struct Env *env_link; // Next free env
  uint32_t env_id; // Unique environment identifier
  uint32_t env_parent_id; // env_id of this env's parent
  enum EnvType env_type; // Indicates special system environments
  unsigned env_status; // Status of the environment
  uint32_t env_runs; // Number of times environment has run

  // Address space
  pde_t *env_pgdir; // Kernel virtual address of page dir
};
```

Here's what the `Env` fields are for:

env_tf
This structure, defined in `inc/trap.h`, holds the saved register values for the environment while that environment is *not running*; i.e., when the kernel or a different environment is running. The kernel saves these when switching from user to kernel mode, so that the environment can later be resumed where it left off.

env_link
This is a link to the next `Env` on the `env_free_list`. `env_free_list` points to the first free environment on the list.

env_id
The kernel stores here a value that uniquely identifies the environment currently using this `Env` structure (i.e., using this particular slot in the `envs` array). After a user environment terminates, the kernel may re-allocate the same `Env` structure to a different environment – but the new environment will have a different `env_id` from the old one even though the new environment is re-using the same slot in the `envs` array.

env_parent_id
The kernel stores here the `env_id` of the environment that created this environment. In this way the environments can form a "family tree," which will be useful for making security decisions about which environments are allowed to do what to whom.

env_type
This is used to distinguish special environments. For most environments, it will be `ENV_TYPE_USER`. We'll introduce a few more types for special system service environments in later labs.

env_status
This variable holds one of the following values:

`ENV_FREE`: Indicates that the `Env` structure is inactive, and therefore on the `env_free_list`.
`ENV_READY`: Indicates that the `Env` structure represents an environment that is waiting to run on the processor.
`ENV_RUNNING`: Indicates that the `Env` structure represents the currently running environment.
`ENV_NOT_RUNNING`: Indicates that the `Env` structure represents a currently active environment, but it is not currently ready to run: for example, because it is waiting for an interprocess communication (IPC) from another environment.
`ENV_DYING`: Indicates that the `Env` structure represents a zombie environment. A zombie environment will be freed the next time it traps to the kernel. We will not use this flag until Lab 4.

env_pgdir
This variable holds the kernel *virtual address* of this environment's page directory.

Like a Unix process, a JOS environment couples the concepts of "thread" and "address space". The *thread* is defined primarily by the saved registers (the `env_tf` field), and the address space is defined by the page directory and page tables pointed to by `env_pgdir`. To run an environment, the kernel must set up the CPU with *both* the saved registers and the appropriate address space.

Our `struct Env` is analogous to `struct proc` in xv6. Both structures hold the environment's (i.e., process's) user-mode register state in a `Trapframe` structure. In JOS, individual environments do not have their own kernel stacks as processes do in xv6. There can be only one `Env` environment active in the kernel at a time, so JOS needs only a *single* kernel stack.

Allocating the Environments Array

In lab 2, you allocated memory in `mem_init()` for the `pages[]` array, which is a table the kernel uses to keep track of which pages are free and which are not. You will now need to modify `mem_init()` further to allocate a similar array of `Env` structures, called `envs`.

Exercise 1. Modify `mem_init()` in `kern/panic.c` to allocate and map the `envs` array. This array consists of exactly `ENV` instances of the `Env` structure allocated much like how you allocated the `pages` array. Also like the `pages` array, the memory backing `envs` should also be mapped user-read-only at `UDRW` (defined in `inc/memlayout.h`) so user processes can read from this array.

You should run your code and make sure `check_kern_pgdir()` succeeds.

Creating and Running Environments

You will now write the code in `kern/env.c` necessary to run a user environment. Because we do not yet have a filesystem, we will set up the kernel to load a static binary image that is *embedded within the kernel itself*. JOS embeds this binary in the kernel as a ELF executable image.

The Lab 3 `makefile` generates a number of binary images in the `obj/user/` directory. If you look at `kern/makefrag`, you will notice some magic that "links" these binaries directly into the kernel executable as if they were `.o` files. The `-b` binary option on the linker command line causes these files to be linked in as "raw" uninterpreted binary files rather than as regular `.o` files produced by the compiler. (As far as the linker is concerned, these outputs do not have to be ELF images at all – they could be anything, such as text files or pictures!) If you look at `obj/kern/kernel.sym` after building the kernel, you will notice that the linker has "magically" produced a number of funny symbols with obscure names like `_binary_obj_user_hello_start`, `_binary_obj_user_hello_end`, and `_binary_obj_user_hello_size`. The linker generates these symbol names by mangling the file names of the binary files; the symbols provide the regular kernel code with a way to reference the embedded binary files.

In `lib64/init.c` in `kern/init.c` you'll see code to run one of these binary images in an environment. However, the critical functions to set up user environments are not complete; you will need to fill them in.

Exercise 2. In the file `env.c`, finish coding the following functions:

```
env_init()
  Initialize all of the Env structures in the envs array and add them to the env_free_list. Also calls env_init_percpu, which configures the segmentation hardware with separate segments for privilege level 0 (kernel) and privilege level 3 (user).
env_setup_vm()
  Allocate a page directory for a new environment and initialize the kernel portion of the new environment's address space.
region_alloc()
  Allocates and maps physical memory for an environment
load_icode()
  You will need to parse an ELF binary image, much like the boot loader already does, and load its contents into the user address space of a new environment.
env_create()
  Allocate an environment with env_alloc and call load_icode to load an ELF binary into it.
env_run()
  Start a given environment running in user mode.
```

As you write these functions, you might find the new `printf` verb `va_usage` – it prints a description corresponding to an error code. For example,

```
panic("env_alloc: %e", r);
will panic with the message "env_alloc: out of memory".
```

Below is a call graph of the code up to the point where the user code is invoked. Make sure you understand the purpose of each step.

- start(kern/entry.S)
- lib64_init(kern/init.c)
 - cons_init
 - mem_init
 - env_init
 - trap_init (still incomplete at this point)
 - env_create
 - env_run
 - env_pop_tf

Once you are done you should compile your kernel and run it under QEMU. If all goes well, your system should enter user space and execute the `hello` binary until it makes a system call with the `int` instruction. At that point there will be trouble, since JOS has not set up the hardware to allow any kind of transition from user space into the kernel. When the CPU discovers that it is not set up to handle this system call interrupt, it will generate a general protection exception, find that it can't handle that, generate a double fault exception, find that it can't handle that either, and finally give up with what's known as a "triple fault". Usually, you would then see the CPU reset and the system reboot. While this is important for legacy applications (see [this blog post](#) for an explanation of why), it's a pain for kernel development, so with the 6.828 patched QEMU you'll instead see a register dump and a "Triple fault." message.

We'll address this problem shortly, but for now we can use the debugger to check that we're entering user mode. Use `make qemu-gdb` and set a GDB breakpoint at `env_pop_tf`, which should be the last function you hit before actually entering user mode. Single step through this function using `si`; the processor should enter user mode after the `iret` instruction. You should then see the first instruction in the user environment's executable, which is the `cmpl` instruction at the label `start` in `lib/entry.S`. Now use `b *0x...` to set a breakpoint at the `int_0x030` in `sys_cpuidt()` in `hello` (see `obj/user/hello.asm` for the user-space address). This `int` is the system call to display a character to the console. If you can't execute as far as the `int`, then something is wrong with your address space setup or program loading code; go back and fix it before continuing.

Handling Interrupts and Exceptions

At this point, the first `int_0x030` system call instruction in user space is a dead end: once the processor gets into user mode, there is no way to get back out. You will now need to implement basic exception and system call handling, so that it is possible for the kernel to recover control of the processor from user-mode code. The first thing you should do is thoroughly familiarize yourself with the x86 interrupt and exception mechanism.

Exercise 3. Read [Chapter 9: Exceptions and Interrupts](#) in the [20386 Programmer's Manual](#) (or Chapter 5 of the [IA-32 Developer's Manual](#)), if you haven't already.

In this lab we generally follow Intel's terminology for interrupts, exceptions, and the like. However, terms such as exception, trap, interrupt, fault and abort have no standard meaning across architectures or operating systems, and are often used without regard to the subtle distinctions between them on a particular architecture such as the x86. When you see these terms outside of this lab, the meanings might be slightly different.

Basics of Protected Control Transfer

Exceptions and interrupts are both "protected control transfers," which cause the processor to switch from user to kernel mode (CPL=0) without giving the user-mode code any opportunity to interfere with the functioning of the kernel or other user-mode code. Single step through this function using `si`; the processor should enter user mode after the `iret` instruction. You should then see the first instruction in the user environment's executable, which is the `cmpl` instruction at the label `start` in `lib/entry.S`. Now use `b *0x...` to set a breakpoint at the `int_0x030` in `sys_cpuidt()` in `hello` (see `obj/user/hello.asm` for the user-space address). This `int` is the system call to display a character to the console. If you can't execute as far as the `int`, then something is wrong with your address space setup or program loading code; go back and fix it before continuing.

In order to ensure that these protected control transfers are actually *synchronized* by the processor's interrupt/exception mechanism is designed so that the code currently running when the interrupt or external device *does not get to choose arbitrarily where the kernel is entered or how*. Instead, the processor ensures that the kernel can be entered only under carefully controlled conditions. On the x86, two mechanisms work together to provide this protection:

- The Interrupt Descriptor Table.** The processor ensures that interrupts and exceptions can only cause the kernel to be entered at a few specific, well-defined entry-points *determined by the kernel itself*, and not by the code running when the interrupt or exception is taken.
- The Task State Segment.** The processor needs a place to save the *old* processor state before the interrupt or exception occurred, such as the original values of `cs` and `esp` before the processor invoked the exception handler, so that the exception handler can later restore that old state and resume the interrupted code from where it left off. But this save area for the old processor state must in turn be protected from unprivileged user-mode code; otherwise buggy or malicious user code could compromise the kernel.

For this reason, when an x86 processor takes an interrupt or trap that causes a privilege level change from user to kernel mode, it also switches to a stack in the kernel's memory. A structure called the *task state segment* (TSS) specifies the segment selector and address where this stack lives. The processor pushes (on this new stack) `ss`, `esp`, `eflags`, `cs`, `esp`, and an optional error code. Then it loads the `cs` and `esp` from the kernel's interrupt descriptor, and sets the `esp` and `ss` to refer to the new stack.

Although the TSS is large and can potentially serve a variety of purposes, JOS only uses it to define the kernel stack that the processor should switch to when it transfers from user to kernel mode. Since "kernel mode" in JOS is privilege level 0 on the x86, the processor uses the `ss0` and `esp0` fields of the TSS to define the kernel stack when entering kernel mode. JOS doesn't use any other TSS fields.

Types of Exceptions and Interrupts

All of the synchronous exceptions that the x86 processor can generate internally use interrupt vectors between 0 and 31, and therefore map to IDT entries 0-31. For example, a page fault always causes an exception through vector 14. Interrupt vectors greater than 31 are only used by *software* interrupts, which can be generated by the `int` instruction, or asynchronous *hardware* interrupts, caused by external devices when they need attention.

In this section we will extend JOS to handle the internally generated x86 exceptions in vectors 0-31. In the next section we will make JOS handle software interrupt vectors 48 (0x30), which JOS (fairly arbitrarily) uses as its system call interrupt vector. In Lab 4 we will extend JOS to handle externally generated hardware interrupts such as the clock interrupt.

An Example

Let's put these pieces together and trace through an example. Let's say the processor is executing code in a user environment and encounters a divide instruction that attempts to divide by zero.

- The processor switches to the stack defined by the `ss0` and `esp0` fields of the TSS, which in JOS will hold the values `0x0` and `KSTACKTOP`, respectively.
- The processor pushes the exception parameters on the kernel stack, starting at address `KSTACKTOP`:

```
0x00000 | old SS      ~ 4
0x00000 | old EFLAGS   ~ 8
0x00000 | old CS       ~ 12
0x00000 | old EIP      ~ 16
0x00000 | old EIP      ~ 20
0x00000 | old EIP      ~ 24 ~ ESP
```

- Because we're handling a divide error, which is interrupt vector 0 on the x86, the processor reads IDT entry 0 and sets `cs:EIP` to point to the handler function described by the entry.
- The handler function takes control and handles the exception, for example by terminating the user environment.

For certain types of x86 exceptions, in addition to the "standard" five words above, the processor pushes onto the stack another word containing an *error code*. The page fault exception, number 14, is an important example. See the 80386 manual to determine for which exception numbers the processor pushes an error code, and what the error code means in that case. When the processor pushes an error code, the stack would look as follows at the beginning of the exception handler when coming from user mode:

```
0x00000 | old SS      ~ 4
0x00000 | old EFLAGS   ~ 8
0x00000 | old CS       ~ 12
0x00000 | old EIP      ~ 16
0x00000 | error code     ~ 20 ~ ESP
```

Nested Exceptions and Interrupts

The processor can take exceptions and interrupts both from kernel and user mode. It is only when entering the kernel from user mode, however, that the x86 processor automatically switches stacks before pushing its old register state to the stack and invoking the appropriate exception handler through the IDT. If the processor is *already* in kernel mode when the interrupt or exception occurs the low 2 bits of the `cs` register are already zero, then the CPU just pushes more values on the same kernel stack. In this way, the kernel can gracefully handle *nested exceptions* caused by code within the kernel itself. This capability is an important tool in implementing protection, as we will see later in the section on system calls.

If the processor is already in kernel mode and takes a nested exception, since it does not need to switch stacks, it does not save the old `ss` or `esp` registers. For exception types that do not push an error code, the kernel stack therefore looks like the following on entry to the exception handler:

```
0x00000 | old CS      ~ 4
0x00000 | old EFLAGS   ~ 8
0x00000 | old EIP      ~ 12
```

For exception types that push an error code, the processor pushes the error code immediately after the old `EIP`, as before.

There is one important caveat to the processor's nested exception capability. If the processor takes an exception while already in kernel mode, and *cannot push its old state onto the kernel stack* for any reason such as lack of stack space, then there is nothing the processor can do to recover, so it simply resets itself. Needless to say, the kernel should be designed so that this can't happen.

Setting Up the IDT

You should now have the basic information you need in order to set up the IDT and handle exceptions in JOS. For now, you will set up the IDT to handle interrupt vectors 0-31 (the processor exceptions). We'll handle system call interrupts later in this lab and add interrupts 32-47 (the device IRQs) in a later lab.

The header files `inc/trap.h` and `kern/trap.h` contain important definitions related to interrupts and exceptions that you will need to become familiar with. The file `kern/trap.h` contains definitions that are strictly private to the kernel, while `inc/trap.h` contains definitions that may also be useful to user-level programs and libraries.

Note: Some of the exceptions in the range 0-31 are defined by Intel to be reserved. Since they will never be generated by the processor, it doesn't really matter how you handle them. Do whatever you think is cleanest.

The overall flow of control that you should achieve is depicted below:

```
IDT      trapentry.S      trap.c
-----> handler1:      trap (struct Trapframe *tf)
// do stuff          { // handle the exception/interrupt
// ...                }
-----> handler2:      // do stuff
// do stuff          call trap
// ...                // ...
.
.
.
-----> handlerX:      // do stuff
// do stuff          call trap
// ...                // ...
```

Each exception or interrupt should have its own handler in `trapentry.S` and `trap_init()` should initialize the IDT with the addresses of these handlers. Each of the handlers should build a `struct Trapframe` (see `inc/trap.h`) on the stack and call `trap()` (in `trap.c`) with a pointer to the `Trapframe`; `trap()` then handles the exception/interrupt or dispatches to a specific handler function.

Exercise 4. Edit `trapentry.S` and `trap.c` and implement the features described above. (The macros `TRAPHANDLER` and `TRAPHANDLER_IRQ0C` in `trapentry.S` should help you, as well as the `T_*` defines in `inc/trap.h`. You will need to add an entry point in `trapentry.S` (using those macros) for each trap defined in `inc/trap.h`, and you'll have to provide `_alltraps` which the `TRAPHANDLER` macros refer to. You will also need to modify `trap_init()` to initialize the `idt` to point to each of these entry points defined in `trapentry.S`; the `SETGATE` macro will be helpful here.)

Your `_alltraps` should:

- push values to make the stack look like a `struct Trapframe`
- load `0x0` into `ss` and `yes`
- push: `0x0` to pass a pointer to the `Trapframe` as an argument to `trap()`
- call `trap()` (can trap over return?)

Consider using the `pushal` instruction; it fits nicely with the layout of the `struct Trapframe`.

Test your trap handling code using some of the test programs in the `user` directory that cause exceptions before making any system calls, such as `user/divzero`. You will have to get `make grade` to succeed on the `divzero`, `softint`, and `badsegment` tests at this point.

Challenge! You probably have a lot of very similar code right now, between the lists of `TRAPHANDLER`s in `trapentry.S` and their installations in `trap.c`. Clean this up. Change the macros in `trapentry.S` to automatically generate a table for `trap.c` to use. Note that you can switch between laying down code and data in the assembler by using the directives `.text` and `.data`.

Questions

Answer the following questions in `lab3-questionary.txt`:

- What is the purpose of having an individual handler function for each exception/interrupt? (i.e., if all exceptions/interrupts were delivered to the same handler, what feature that exists in the current implementation could not be provided?)
- Did you have to do anything to make the `user/softint` program behave correctly? The grade script expects it to produce a general protection fault (trap 13), but `softint.c` says `int_0x14`. Why should this produce interrupt vector 13? What happens if the kernel actually allows `softint's int_0x14` instruction to invoke the kernel's page fault handler (which is interrupt vector 14)?

This concludes Part A of the lab.

Part B: Page Faults, Exceptions, and System Calls

Now that your kernel has basic exception handling capabilities, you will refine it to provide important operating system primitives that depend on exception handling.

Handling Page Faults

The page fault exception, interrupt vector 14 (`#PF`), is a particularly important one that we will exercise heavily throughout this lab and the next. When the processor takes a page fault, it stores the linear (i.e., virtual) address that caused the fault in the processor control register, `cr2`. In `trap.c` we have provided the beginnings of a special function, `page_fault_handler()`, to handle page fault exceptions.

Exercise 5. Modify `trap_dispatch()` to dispatch page fault exceptions to `page_fault_handler()`. You should now be able to get `make grade` to succeed on the `faultread`, `faultwrite`, `faultexec`, and `faultexec` kernel tests. If any of them don't work, figure out why and fix them. Remember that you can boot JOS into a particular user program using `make user.a` or `make user-a.asm`.

You will further refine the kernel's page fault handling below, as you implement system calls.

The Breakpoint Exception

The breakpoint exception, interrupt vector 3 (`#BP`), is normally used to allow debuggers to insert breakpoints in a program's code by temporarily replacing the relevant program instruction with the special 1-byte `int3` software interrupt instruction. In JOS we will abuse this exception slightly by turning it into a primitive pseudo-system call that any user environment can use to invoke the JOS kernel monitor. This usage is actually somewhat appropriate if we think of the JOS kernel monitor as a primitive debugger. The user-mode implementation of `panic()` in `lib/panic.c`, for example, performs an `int3` after displaying its panic message.

Exercise 6. Modify `trap_dispatch()` to make breakpoint exceptions invoke the kernel monitor. You should now be able to get `make grade` to succeed on the `breakpoint` test.

Challenge! Modify the JOS kernel monitor so that you can "continue" execution from the current location (e.g., after the `int3`), if the kernel monitor was invoked using a page fault exception, and so that you can single-step one instruction at a time. You will need to understand certain bits of the `eflags` register in order to implement single-stepping.

Optional: If you're feeling really adventurous, find some x86 disassembler source code – e.g., by ripping it out of QEMU, or out of GNU binutils, or just write it yourself – and extend the JOS kernel monitor to be able to disassemble and display instructions as you go, by stepping through them. Combined with the symbol table loading from lab 2, this is the stuff of which real kernel debuggers are made.

Questions
3. The breakpoint test case will either generate a break point exception or a general protection fault depending on how you initialized the global pointer entry in the IDT (i.e., you call to `SETGATE` from `trap_init`). Why? How do you need to set it up in order to get the breakpoint exception to work as specified above and what incorrect setup would cause it to trigger a general protection fault?

4. What do you think is the point of these mechanisms, particularly in light of what the `user/softint` test program does?

System calls

User processes ask the kernel to do things for them by invoking system calls. When the user process invokes a system call, the processor enters kernel mode, the processor and the kernel cooperate to save the user process's state, the kernel executes appropriate code in order to carry out the system call, and then resumes the user process. The exact details of how the user process gets the kernel's attention and how it specifies which call it wants to execute vary from system to system.

In the JOS kernel, we will use the `int` instruction, which causes a processor interrupt. In particular, we will use `int_0x030` as the system call interrupt. We have defined the constant `T_SYSCALL` in `48 (0x30)` for you. You will have to set up the interrupt descriptor to allow user processes to cause that interrupt. Note that interrupt 0x30 cannot be generated by hardware, so there is no ambiguity caused by allowing user code to generate it.

The application will pass the system call number and the system call arguments in registers. This way, the kernel won't need to grub around in the user environment's stack or instruction stream. The system call number will go in `eax`, and the arguments (up to five of them) will go in `edx`, `ecx`, `ebx`, `edi`, and `esi`, respectively. The kernel passes the return value back in `eax`. The assembly code to invoke a system call has been written for you, in `syscall1()` in `lib/syscall.c`. You should read through it and make sure you understand what is going on.

Exercise 7. Add a handler in the kernel for interrupt vector `T_SY`